

# STAT 5630, Fall 2019

## Clustering

---

Xiwei Tang, Ph.D. <[xt4yj@virginia.edu](mailto:xt4yj@virginia.edu)>

University of Virginia  
October 29, 2019

# Unsupervised Learning

---

# Unsupervised Learning

- Learn without the response variable  $Y$ , only  $\{x_i\}_{i=1}^n$
- Goal: learn patterns in  $X$
- Examples
  - Estimate the **density**, **covariance**, **graph (network)**, etc. of  $X$  — could be difficult high-dimensional problems
  - **Dimension reduction**: identify low-dimensional manifolds within the feature space  $\mathcal{X}$  that represent high data density.
  - **Cluster analysis**: try to identify multiple regions of the feature space that contains modes of density.
- Oftentimes, there is no clear measure of success for unsupervised learning.

# Cluster Analysis

---

# Cluster Analysis

- Group the dataset into subsets so that those within each subset are more closely related (similar) to each other than those objects assigned to other subsets. Each subset is called a cluster
- Flat clustering vs. hierarchical clustering: flat clustering divides the dataset into  $k$  clusters; and hierarchical clustering is to arrange the clusters into a natural hierarchy.
- Clustering results are crucially dependent on the measure of similarity (or distance) between the “points” to be clustered.

# Distance Matrices

- **Dissimilarity Matrix** is distance matrices that gives small values for a pair of similar (closer) subjects.
- The distance between two vectors  $\mathbf{x}_i$  and  $\mathbf{x}_{i'}$  is given by

$$D(\mathbf{x}_i, \mathbf{x}_{i'}) = \sum_{j=1}^p d_j(x_{ij}, x_{i'j})$$

- The most commonly used measurement (for continuous features) is the **Euclidian distance**:

$$d_j(x_{ij}, x_{i'j}) = (x_{ij} - x_{i'j})^2,$$

hence  $D(\mathbf{x}_i, \mathbf{x}_{i'}) = \|\mathbf{x}_i - \mathbf{x}_{i'}\|_2^2$ .

# Distance Measures

- The Euclidian distance can be related to the correlation (angle) between the two vectors  $\mathbf{x}_i$  and  $\mathbf{x}_{i'}$ .
- In high-dimensional settings, we sometimes center and standardize (to mean 0 and sd 1) the vectors  $\mathbf{x}_i$  for each subject.
- Then the correlation defined as

$$\rho(\mathbf{x}_i, \mathbf{x}_{i'}) = \frac{\sum_j (x_{ij} - \bar{x}_i)(x_{i'j} - \bar{x}_{i'})}{\sqrt{\sum_j (x_{ij} - \bar{x}_i)^2 \sum_j (x_{i'j} - \bar{x}_{i'})^2}}$$

is related to the distance measure through the relationship

$$D(x_{ij}, x_{i'j}) = 2(1 - \rho(\mathbf{x}_i, \mathbf{x}_{i'}))$$

- **Note:** this standardization is only useful in high-dimensional settings.

# Clustering Algorithms

- We want to form  $K$  clusters, indexed by  $k \in \{1, \dots, K\}$ .
- Let  $C(i)$  : denote the cluster index assigned to the  $i$ th observation.
- Goal: search for the function  $C : \{1, \dots, n\} \rightarrow \{1, \dots, K\}$  to minimize the overall **within cluster distance**:

$$W(C) = \frac{1}{2} \sum_{k=1}^K \sum_{C(i), C(i')=k} d(x_i, x_{i'}).$$



- Similarly, we could maximize the **between cluster distance**

$$B(C) = \frac{1}{2} \sum_{k=1}^K \sum_{C(i)=k} \sum_{C(i') \neq k} d_{ii'}$$

- However, the **total distance** can be broke down into

$$\begin{aligned} T &= \frac{1}{2} \sum_{i=1}^n \sum_{i'=1}^n d_{ii'} = \frac{1}{2} \sum_{k=1}^K \sum_{C(i)=k} \left[ \sum_{C(i')=k} d_{ii'} + \sum_{C(i') \neq k} d_{ii'} \right] \\ &= W(C) + B(C) \end{aligned}$$

- The total distance is fixed for a given set of data.

$$\text{Minimizing } W(C) \iff \text{maximizing } B(C)$$

# Clustering Algorithms

- Given a specific distance measure  $d(\cdot, \cdot)$ , several algorithms can be used to find the clusters
  - Combinatorial algorithms
  - $K$ -means clustering
  - Hierarchical clustering

- For **small**  $n$  and  $K$ , we could minimize  $W$  by brute-force search.
- However, this is not feasible for large  $n$  and  $K$ , since the number of distinct assignments can be extremely large:

$$S(n, K) = \frac{1}{K!} \sum_{k=1}^K (-1)^{K-k} \binom{K}{k} k^n$$

- For example  $S(10, 4) = 34,105$ ;  $S(19, 4) \approx 10^{10}$ .
- It calls for more efficient algorithms: may not be optimal but a reasonably good suboptimal partition.

- Idea: an iterative greedy descent algorithm
  - **Initialization**: A random partition of  $K$  clusters is specified.
  - **Iterative step**: Update the cluster registration towards a direction that the within cluster distance is reduced.
  - **Stopping rule**: When no improvement can be reached, terminate the algorithm
- Convergence is guaranteed, but not necessarily to global optima.

- Consider an enlarged optimization problem:

$$\min_{C, \{m_k\}_{k=1}^K} \sum_{k=1}^K \sum_{C(i)=k} \|x_i - m_k\|^2$$

- This problem is NP-hard for  $\geq 2$  dimensions.
- Instead, consider an algorithm that alternatively update the two components:
  - $C$ , the cluster assignments
  - $\{m_k\}_{k=1}^K$ : the cluster means
- We will do an iterative update by:
  - Fixing  $C$ , find the best  $\{m_k\}_{k=1}^K$
  - Fixing  $\{m_k\}_{k=1}^K$ , find the best  $C$

- Fixing  $C$ , we know the cluster label of each subject. For any set  $\{i : C(i) = k\}$ , finding the mean is

$$m_k = \arg \min_m \sum_{C(i)=k} \|x_i - m\|^2.$$

This is simply finding the mean of all observations with  $C(i) = k$ .

- Fixing the cluster means  $\{m_k\}_{k=1}^K$ , to find the new cluster assignments, we simply recalculate the distance from an observation to each of the cluster mean.

$$C(i) = \arg \min_k d(x_i, \bar{x}_k)$$

- A  $K$ -means Clustering algorithm:
  - 1) Random split the dataset into  $K$  different subsets. Assign each subsets a cluster label. Then iterate between 2) and 3).
  - 2) Given cluster assignments  $C$ , calculate the cluster mean vectors  $m_1, \dots, m_K$ .
  - 3) Given a current set of means  $\{m_1, \dots, m_K\}$ , assign each observation to the closest current cluster mean.
- **Note:** We usually initiate the cluster labels randomly. However, this algorithm does not guarantee to converge to the global minimizer.
- The algorithm still has a **descent property**, which leads to a local minimizer.

- $K$ -medoids is an alternative version of  $K$ -means:
- Replace the second step by searching for the **observation** that minimizes the distance to all others in the cluster

$$i_k^* = \arg \min_{i: C(i)=k} \sum_{C(i')=k} D(x_i, x_{i'})$$

- Use  $x_{i_k^*}$  as the “center” of cluster  $k$ .




- An application of  $K$ -means: image segmentation (example taken from Prof. Domeniconi, GMU)
- **Goal:** partition an image into regions with homogeneous visual appearance (which could correspond to objects or parts of objects). This could save storage space.
- **Image representation:** each pixel is represented as a three dimensional point in RGB space, where
  - R = intensity of red
  - G = intensity of green
  - B = intensity of blue

# Image segmentation

Original image



pixel




# Image segmentation

Original image



Mean (R,G,B)



# Image segmentation

Original image



$K = 2$



$K = 3$




$K = 10$




# Hierarchical Clustering

- Choosing the number of clusters  $K$  can be difficult
- A hierarchical representation which
  - at the lowest level, each cluster contains a single observation.
  - at the highest level there is only one cluster containing all observations.
- Use dendrogram to display the clustering result.
- Two paradigms: agglomerative (bottom-up) and divisive (top-down).
- An example taken from Prof. Ullman, Stanford


# Example: Hierarchical Clustering



# Example: Hierarchical Clustering




# Example: Hierarchical Clustering






# Example: Hierarchical Clustering



after several steps

## Example: Hierarchical Clustering



A tree representation

# Agglomerative Clustering (bottom-up)

- Begin with every observation representing a singleton cluster.
- At each step, merge two “closest” clusters into one cluster and reduce the number of clusters by one.
- Need a measure of dissimilarity between two clusters.
- Dissimilarity between two clusters  $G$  and  $H$ :  $d(G, H)$ . Different choices:
  - Single linkage: the closest pair  $d(G, H) = \min_{i \in G, i' \in H} d_{ii'}$
  - Complete linkage: the furthest pair  $d(G, H) = \max_{i \in G, i' \in H} d_{ii'}$
  - Group Average: average dissimilarity
$$d(G, H) = \frac{1}{n_G n_H} \sum_{i \in G} \sum_{i' \in H} d_{ii'}$$

- How many clusters we should choose?
- Sparse clustering for high-dimensional data?
- R implementation (base package):
  - `kmeans`
  - `hclust`