

# STAT 5630, Fall 2019

## Linear Models for Regression

---

Xiwei Tang, Ph.D. <[xt4yj@virginia.edu](mailto:xt4yj@virginia.edu)>

University of Virginia  
September 10, 2019

- Linear Models for Regression
- Model Selection Criteria

# Linear Models for Regression

---

# Regression Models

- Observe a collection of i.i.d. **training data**

$$\mathcal{D}_n = \{\mathbf{x}_i, y_i\}_{i=1}^n$$

where each  $\mathbf{x}_i$  is a  $p$  dimensional vector (**prediction variables**, covariates, features, inputs), i.e.

$$\mathbf{x}_i = (x_{i1}, \dots, x_{ip})^\top$$

and  $y_i \in \mathbb{R}$  is a **continuous response** (outcome, output).

- Denote  $\mathbf{X}_{\cdot j}$  is a  $n$  dimensional vector of the  $j$ th feature, i.e.

$$\mathbf{X}_{\cdot j} = (x_{1j}, x_{2j}, \dots, x_{nj})^\top$$

- The design matrix  $\mathbf{X}$  is  $n \times p$  dimensional,

$$\mathbf{X} = (\mathbf{X}_{\cdot 1}, \mathbf{X}_{\cdot 2}, \dots, \mathbf{X}_{\cdot p})$$

# Loss and Risk functions

- To estimate  $f(X)$ , need to define a criterion for model fitting.
- A **loss** function  $L$  measures discrepancies between  $Y$  and  $f(X)$ .
- A commonly used loss function is the **squared error loss**:

$$L(Y, f(X)) = (Y - f(X))^2.$$

- **Risk** is the expected loss over the entire population

$$R(f) = E [L(Y, f(X))] = E [(Y - f(X))^2].$$

# Minimizing the Empirical Risk


- With the training data  $\mathcal{D}_n$ , we can solve for a  $f(x)$  such that the **empirical risk** is minimized, i.e., we replace the expectation with the average over  $n$  training samples:

$$\frac{1}{n} \sum_{i=1}^n L(y_i - f(x_i)).$$

- Using the squared error, we search for a function  $\hat{f}$  to minimize the empirical risk on the training dataset

$$\hat{f} = \arg \min_{f \in \mathcal{F}} \sum_{i=1}^n (y_i - f(x_i))^2.$$

- $\mathcal{F}$  is a space of models that we consider.



**FIGURE 2.10.** *Least squares fitting of a function of two inputs. The parameters of  $f_{\theta}(x)$  are chosen so as to minimize the sum-of-squared vertical errors.*

# Linear Regression

- To estimate  $f$ , we pose some **restrictions/structures**
- A **linear regression** model describes the dependence between  $X$  and  $Y$  by

$$\begin{aligned} Y &= X^T \beta + \epsilon \\ &= \beta_1 X_1 + \cdots + \beta_p X_p + \epsilon \end{aligned}$$

where we can set  $X_1 = 1$  as the intercept if necessary, and  $\epsilon$  is an independent error term.

- Given the training data  $\mathcal{D}_n$ , we express the regression model in the **matrix form**

$$\mathbf{y}_{n \times 1} = \mathbf{X}_{n \times p} \boldsymbol{\beta}_{p \times 1} + \mathbf{e}_{n \times 1}$$

where  $\mathbf{X}_{n \times p}$  is the **design matrix** with each row representing an input vector from one subject.




- Now, the estimation of  $f$  boils down to the estimating of  $\beta$
- By our previous definition of the empirical risk, we try to solve for  $\beta$  that minimizing the residual sum of squares (RSS)

$$\begin{aligned}\text{RSS} &= \sum_{i=1}^n \left( y_i - x_{i1}\beta_1 - \cdots, x_{ip}\beta_p \right)^2 \\ &= \|\mathbf{y} - \mathbf{X}\beta\|^2 \\ &= (\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta)\end{aligned}$$

- The ordinary least squares estimator (OLS) is

$$\hat{\beta} = \arg \min_{\beta} (\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta)$$

# Linear Regression



**FIGURE 3.1.** *Linear least squares fitting with  $X \in \mathbb{R}^2$ . We seek the linear function of  $X$  that minimizes the sum of squared residuals from  $Y$ .*

- To estimate  $\beta$ , we set the derivative equal to 0

$$\begin{aligned}\frac{\partial \text{RSS}}{\partial \beta} &= -2\mathbf{X}^\top (\mathbf{y} - \mathbf{X}\beta) = 0 \\ \implies \mathbf{X}^\top \mathbf{y} &= \mathbf{X}^\top \mathbf{X} \beta\end{aligned}$$

which is commonly known as the **normal equation**.

- We then have, if  $\mathbf{X}^\top \mathbf{X}$  is invertible,

$$\hat{\beta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}.$$

- $\mathbf{X}$  full rank  $\iff \mathbf{X}^\top \mathbf{X}$  invertible

- The fitted values (prediction at the observed data points) are

$$\hat{\mathbf{y}} = \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y} \doteq \mathbf{H}_{n \times n} \mathbf{y}$$

- $\mathbf{H}$  (“hat matrix”) is a **project** matrix
  - symmetric:  $\mathbf{H}^\top = \mathbf{H}$
  - idempotent:  $\mathbf{H}\mathbf{H} = \mathbf{H}$
- The **residual**  $\mathbf{r}_{n \times 1} = \hat{\mathbf{e}} = \mathbf{y} - \hat{\mathbf{y}} = (\mathbf{I} - \mathbf{H})\mathbf{y}$
- $\mathbf{r}$  can be used to estimate the error variance


$$\hat{\sigma}^2 = \frac{1}{n-p} \sum_{i=1}^n r_i^2 = \frac{\text{RSS}}{n-p}$$

# Vector Space Interpretation

- The **essence of LS** is to decompose the data vector  $y$  into two orthogonal vectors

$$y = \hat{y} + r$$

- Note that the normal equations implies that  $r$  is orthogonal to each column of  $X$ , i.e.,  $X^T r = 0$



# Statistical Properties of $\hat{\beta}$

- We often assume that the samples are generated from the model

$$Y = X^T \beta + \epsilon,$$

where the errors  $\epsilon_i$  are i.i.d. with  $E(\epsilon_i) = 0$  and  $\text{Var}(\epsilon_i) = \sigma^2$

- Then  $\hat{\beta}$  is unbiased:  $E(\hat{\beta}) = \beta$
- Variance-covariance

$$\text{Var}(\hat{\beta}) = (\mathbf{X}^T \mathbf{X})^{-1} \sigma^2$$

- By the Gauss-Markov Theorem,  $\hat{\beta}$  is the **best linear unbiased estimator** (BLUE)
- But it is **not necessarily** the minimum variance unbiased estimator (MVUE), unless the errors are generated from the Gaussian distribution

# Two Computational Algorithms

- Consider the normal equation:

$$\mathbf{X}^T \mathbf{X} \boldsymbol{\beta} = \mathbf{X}^T \mathbf{y}$$

We would like to avoid computing  $(\mathbf{X}^T \mathbf{X})^{-1}$  directly.

- **QR-decomposition** of  $\mathbf{X}^T \mathbf{X}$ 
  - $\mathbf{X}^T \mathbf{X} = \mathbf{Q} \mathbf{R}$  where  $\mathbf{Q}$  is orthonormal and  $\mathbf{R}$  is upper triangular
- **Cholesky decomposition** of  $\mathbf{X}^T \mathbf{X}$ 
  - $\mathbf{X}^T \mathbf{X} = \mathbf{R} \mathbf{R}^T$ , where  $\mathbf{R}$  is lower triangular

# QR Decomposition

- The QR-decomposition is related to the Gram-Schmidt procedure (see appendix):

$$\mathbf{X}^T \mathbf{X} = \mathbf{Z} \mathbf{\Gamma}$$

where  $\mathbf{\Gamma}$  is an upper triangular matrix and  $\mathbf{Z} = [\mathbf{z}_1, \dots, \mathbf{z}_p]$  has orthogonal columns

- Standardize  $\mathbf{Z}$  using  $\mathbf{D} = \text{diag}\{\|\mathbf{z}_1\|, \dots, \|\mathbf{z}_p\|\}$ ,

$$\mathbf{X}^T \mathbf{X} = \mathbf{Z} \mathbf{\Gamma} = \mathbf{Z} \mathbf{D}^{-1} \mathbf{D} \mathbf{\Gamma} = \mathbf{Q} \mathbf{R}$$

with  $\mathbf{Q} = \mathbf{Z} \mathbf{D}^{-1}$  and  $\mathbf{R} = \mathbf{D} \mathbf{\Gamma}$

- The normal equation reduces to solving the linear system

$$\mathbf{R} \boldsymbol{\beta} = \mathbf{Q}^T \mathbf{X}^T \mathbf{y}$$



# Cholesky Decomposition

- For any positive definite square matrix  $\mathbf{A}$ , we have

$$\mathbf{A} = \mathbf{R}\mathbf{R}^T$$

where  $\mathbf{R}$  is a lower triangular matrix of full rank

- Factoring  $\mathbf{X}^T\mathbf{X} = \mathbf{R}\mathbf{R}^T$
- Solve the triangular system  $\mathbf{R}\mathbf{w} = \mathbf{X}^T\mathbf{y}$
- Solve the triangular system  $\mathbf{R}^T\boldsymbol{\beta} = \mathbf{w}$  for  $\boldsymbol{\beta}$

# Variable Selection

---

- In many applications nowadays, we have many explanatory variables, i.e.,  $p$  is large or even  $p \gg n$ .
  - There are more than 20,000 human protein-coding genes
  - About 10 million single nucleotide polymorphisms (SNPs)
  - Number of subjects,  $n$ , is usually in hundreds or thousands
- In some applications, the key question is to identify a subset of  $X$  variables that are most relevant to  $Y$
- Inadequate to look at marginal effects (P-values)

# Training vs. Testing error

- Training data  $\mathcal{D}_n = \{x_i, y_i\}_{i=1}^n$
- Suppose  $\{x_i, y_i^*\}_{i=1}^n$  is an independent (imaginary) testing dataset collected at the same location  $x_i$ 's (aka, in-sample prediction)
- Assume that the data are indeed from a linear model

$$\begin{aligned} y &= \mu + e = X\beta + e \\ y^* &= \mu + e^* = X\beta + e^* \end{aligned}$$

where both  $y$  and  $y^*$  are  $n \times 1$  response vectors,  $e$  and  $e^*$  are i.i.d. error terms with mean 0 and variance  $\sigma^2$ .

- The true model is indeed linear

# Training vs. Testing error

$$\begin{aligned}E[\text{Test Err}] &= E\|\mathbf{y}^* - \mathbf{X}\hat{\boldsymbol{\beta}}\|^2 \\&= E\|(\mathbf{y}^* - \mathbf{X}\boldsymbol{\beta}) + (\mathbf{X}\boldsymbol{\beta} - \mathbf{X}\hat{\boldsymbol{\beta}})\|^2 \\&= E\|\mathbf{y}^* - \boldsymbol{\mu}\|^2 + E\|\mathbf{X}(\hat{\boldsymbol{\beta}} - \boldsymbol{\beta})\|^2 \\&= E\|\mathbf{e}^*\|^2 + \text{Trace}(\mathbf{X}^\top \mathbf{X} \text{Cov}(\hat{\boldsymbol{\beta}})) \\&= n\sigma^2 + p\sigma^2\end{aligned}$$

$$\begin{aligned}E[\text{Train Err}] &= E\|\mathbf{y} - \hat{\mathbf{y}}\|^2 = E\|(\mathbf{I} - \mathbf{H})\mathbf{y}\|^2 \\&= E\|(\mathbf{I} - \mathbf{H})\mathbf{e}\|^2 \\&= \text{Trace}((\mathbf{I} - \mathbf{H})^\top (\mathbf{I} - \mathbf{H}) \text{Cov}(\mathbf{e})) \\&= (n - p)\sigma^2\end{aligned}$$

So the **testing error** increase with  $p$  and **training error** decreases with  $p$ . When  $p$  gets large, this is a big trouble...

- Hence, it is necessary to select a set of relevant variables, especially when  $p$  is large.
- Variable selection may improve
  - Prediction accuracy
  - Interpretability
- This is a difficult task
  - No natural ordering of importance for the variables
  - The role of a variable needs be measured conditioning on others, high correlation causes trouble
  - It is essential to check all possible combinations, however, this may be computationally expansive

- Model selection is usually done in the following way
  - 1 Give each model a score
  - 2 Design an algorithm to find the model with the best (smallest) score
- The score of a model fitting takes the the form

$$\text{Goodness-of-fit} + \text{Complexity-Penalty}$$

- The first term will decrease as the model gets more complicated (recall 1NN)
- From last week's lecture, the second term increases with the number of predictor variables, which prefers “smaller” model

- Popular choices of scores:
  - Mallows'  $C_p$  (Mallows 1973):  $\text{RSS} + 2\hat{\sigma}_{\text{full}}^2 \cdot p$
  - AIC (Akaike 1970):  $-2 \text{ Log-likelihood} + 2 \cdot p$
  - BIC (Schwarz, 1978):  $-2 \text{ Log-likelihood} + \log n \cdot p$
- When  $n$  is large, adding an additional predictor costs a lot more in BIC than AIC (or  $C_p$ ). So AIC tends to pick a larger model than BIC.
- $C_p$  performs similarly to AIC.



# Stepwise Regression

- **Greedy algorithms**: fast, but only return a local optimal solution (which might be good enough in practice).
  - **Backward**: start with the full model and sequentially delete predictors until the score does not improve.
  - **Forward**: start with the null model and sequentially add predictors until the score does not improve.
  - **Stepwise**: consider both deleting and adding one predictor at each stage.

# Best Subset Selection

- Best subset selection is a **level-wise search algorithm**, which returns the **global optimal** solution for a given model size.
- Only feasible for  $p$  not very large ( $< 50$ )
- Algorithm:
  - 1 For each  $k = 1, \dots, p$ , check  $2^k$  possible combinations, and find the model with smallest RSS
    - The penalty term is the same for models with the same size
  - 2 To choose the best  $k$ , use model selection criteria
- **Note:** if  $\text{RSS}(X_1, X_2) < \text{RSS}(X_3, X_4, X_5, X_6)$  then we do not need to visit any size 2 or 3 sub-models of  $(X_3, X_4, X_5, X_6)$ , which can be **leaped** over.
- Implemented in **R** contributed package “*leaps*”, using the leaps and bounds algorithm (Furnival and Wilson, 1974)

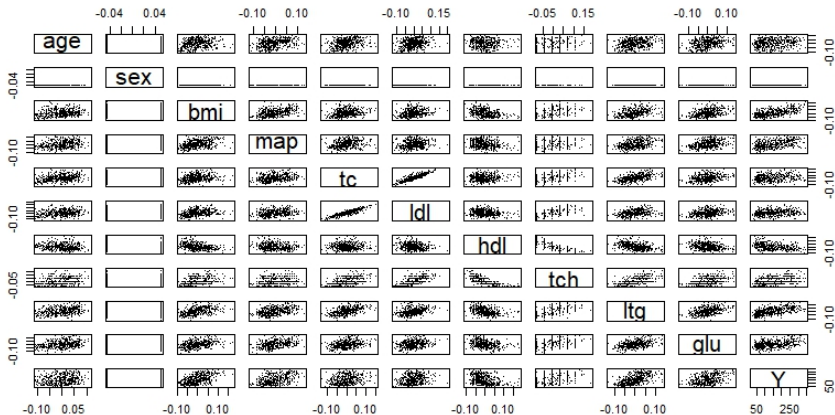
# Diabetes Data Analysis

- The Diabetes Data (Efron et al, 2004) contains ten baseline variables from 442 subjects: age, sex, body mass index, average blood pressure, and six blood serum measurements
- The goal is to model a quantitative measure of disease progression one year after baseline
- Data can be loaded from the [R](#) package “[lars](#)”
- We perform model selections on this dataset (see [R](#) code from course material)


# Diabetes Data

```
1 # Get the Diabetes Data from the package "lars"
2 install.packages("lars")
3 library(lars)
4 data(diabetes)
5 diab = data.frame(cbind(diabetes$x, "Y" = diabetes$y))
6
7 #some data summarization
8 pairs(diab, pch=".") # produce pair-wise scatter plots. Caution
   : this is a big figure.
9
10 # a fancier plot, requires another package
11 install.packages("PerformanceAnalytics")
12 library(PerformanceAnalytics)
13 suppressWarnings(chart.Correlation(diab, col = "purple", pch = "
   *"))
```

# Diabetes Data



# Diabetes Data



# Diabetes Data

```
1 # Run linear regression
2 lmfit=lm(Y~., data=diab)
3 names(lmfit) # What have been returned by "lm"?
4 lmfit$coef   # 11 regression coefficients
5
6 summary(lmfit)
7 round(summary(lmfit)$coef,3) # coefficients and the p-values
8
9 >
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	152.133	2.576	59.061	0.000
age	-10.012	59.749	-0.168	0.867
sex	-239.819	61.222	-3.917	0.000
bmi	519.840	66.534	7.813	0.000
map	324.390	65.422	4.958	0.000
tc	-792.184	416.684	-1.901	0.058
ldl	476.746	339.035	1.406	0.160
hdl	101.045	212.533	0.475	0.635
tch	177.064	161.476	1.097	0.273
ltg	751.279	171.902	4.370	0.000
glu	67.625	65.984	1.025	0.306

# Diabetes Data

```
1 # Step-wise Variable Selection
2 ?step # or help(step)
3 step(lmfit, direction="both") # AIC
4 step(lmfit, direction="both", trace=0) # do not print
   intermediate results
5
6 step(lmfit, direction="backward")
7 step(lm(Y~1, data=diab), scope=list(upper=lmfit, lower=~1),
   direction="forward")
8
9 #n = nrow(diab)
10 step(lmfit, direction="both", k=log(n)) # BIC (the default
   value for k=2, which corresponds to AIC)
11
12 >Call:
13 lm(formula = Y ~ sex + bmi + map + tc + ldl + ltg, data = diab)
14 Coefficients:
15 (Intercept)    sex      bmi      map      tc      ldl      ltg
16 152.1        -226.5   529.9   327.2   -757.9   538.6   804.2
```



# Diabetes Data

```
1 # Best subset model selection (Cp, AIC, and BIC): leaps
2 install.packages("leaps")
3 library(leaps)
4 RSSleaps=regsubsets(as.matrix(diab[, -11]), diab[, 11], nvmax=10)
5 summary(RSSleaps, matrix=T)
6 >1 subsets of each size up to 10
7 Selection Algorithm: exhaustive
8           age sex bmi map tc  ldl hdl tch ltg glu
9 1  ( 1 )    ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** **
10 2  ( 1 )    ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** **
11 3  ( 1 )    ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** **
12 4  ( 1 )    ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** ** **
13 5  ( 1 )    ** ** ** * ** * ** * ** * ** * ** * **
14 6  ( 1 )    ** ** ** * ** * ** * ** * ** * ** * **
15 7  ( 1 )    ** ** ** * ** * ** * ** * ** * ** * **
16 8  ( 1 )    ** ** ** * ** * ** * ** * ** * ** * **
17 9  ( 1 )    ** ** ** * ** * ** * ** * ** * ** * **
18 10 ( 1 )    ** * ** * ** * ** * ** * ** * ** * **
```

# Appendix

---

# Gram-Schmidt Procedure (Successive Orthogonalization)

- In general, to solve for OLS fit (with an intercept), the Gram-Schmidt procedure:
  - 1 Initialize  $\mathbf{z}_0 = \mathbf{x}_0 = \mathbf{1}$
  - 2 For  $j = 1, \dots, p$ , project  $\mathbf{x}_j$  on orthogonal basis  $\{\mathbf{z}_0, \dots, \mathbf{z}_{j-1}\}$ , and let the residual vector be  $\mathbf{z}_j$ . More precisely, let

$$\hat{\gamma}_{kj} = \frac{\langle \mathbf{x}_j, \mathbf{z}_k \rangle}{\langle \mathbf{z}_k, \mathbf{z}_k \rangle} \quad \text{for } k = 0, \dots, j-1,$$

$$\text{and } \mathbf{z}_j = \mathbf{x}_j - \sum_{k=0}^{j-1} \hat{\gamma}_{kj} \mathbf{z}_k$$

- 3 Regress  $\mathbf{y}$  on the residual  $\mathbf{z}_p$  (the last covariate) to get

$$\hat{\beta}_p = \frac{\langle \mathbf{y}, \mathbf{z}_p \rangle}{\langle \mathbf{z}_p, \mathbf{z}_p \rangle}$$

- 4 Compute  $\beta_j$  successively for  $j = p-1, \dots, 0$ .

# Ozone data example for Gram-Schmidt

```
1 > library(ElemStatLearn) # for ozone data
2 > # fit linear regression to the full data
3 > fullmodel = lm(ozone ~ radiation + temperature + wind, data=
  ozone)
4 > round(summary(fullmodel)$coef, dig=3)
5
6           Estimate Std. Error t value Pr(>|t|)
7 (Intercept)  -64.232    23.042  -2.788   0.006
8 radiation      0.060     0.023   2.580   0.011
9 temperature    1.651     0.253   6.516   0.000
10 wind          -3.338     0.654  -5.105   0.000
11
12 > # get the residual of wind after projection on others
13 > wind.res = lm(wind ~ radiation + temperature, data=ozone)$res
14 > # get the coefficient of wind
15 > parmodel=lm(ozone ~ -1 + wind.res, data=ozone)
16 > round(summary(parmodel)$coef, dig=3)
17
18           Estimate Std. Error t value Pr(>|t|)
19 wind.res      -3.338     1.631  -2.046   0.043
```

- Recall our previous analysis of the training and testing errors with  $\mathbf{y}$  and  $\mathbf{y}^*$
- Now, let's assume that the model is not necessarily a linear model, i.e.,

$$\mathbf{y} = \boldsymbol{\mu} + \mathbf{e}$$

$$\mathbf{y}^* = \boldsymbol{\mu} + \mathbf{e}^*$$

We assume mean 0 and variance  $\sigma^2$  for the two error vectors, but we don't have  $\boldsymbol{\mu} = \mathbf{X}\boldsymbol{\beta}$ . However, we still perform linear regression regardless. This will introduce bias of the estimations.

$$\begin{aligned} E[\text{Test Err}] &= E\|\mathbf{y}^* - \mathbf{X}\hat{\boldsymbol{\beta}}\|^2 = \|\mathbf{y}^* - \mathbf{H}\mathbf{y}\|^2 \\ &= E\|(\mathbf{y}^* - \boldsymbol{\mu}) + (\boldsymbol{\mu} - \mathbf{H}\boldsymbol{\mu}) + (\mathbf{H}\boldsymbol{\mu} - \mathbf{H}\mathbf{y})\|^2 \\ &= E\|\mathbf{y}^* - \boldsymbol{\mu}\|^2 + E\|\boldsymbol{\mu} - \mathbf{H}\boldsymbol{\mu}\|^2 + E\|\mathbf{H}\boldsymbol{\mu} - \mathbf{H}\mathbf{y}\|^2 \\ &= E\|\mathbf{e}^*\|^2 + E\|\boldsymbol{\mu} - \mathbf{H}\boldsymbol{\mu}\|^2 + E\|\mathbf{H}\mathbf{e}\|^2 \\ &= n\sigma^2 + \text{Bias}^2 + p\sigma^2 \end{aligned}$$

$$\begin{aligned} E[\text{Train Err}] &= E\|\mathbf{y} - \hat{\mathbf{y}}\|^2 = E\|(\mathbf{I} - \mathbf{H})\boldsymbol{\mu} + (\mathbf{I} - \mathbf{H})\mathbf{e}\|^2 \\ &= E\|(\mathbf{I} - \mathbf{H})\boldsymbol{\mu}\|^2 + E\|(\mathbf{I} - \mathbf{H})\mathbf{e}\|^2 \\ &= \text{Bias}^2 + (n - p)\sigma^2 \end{aligned}$$

Hence, **Test Err** is approximately **Train Err** +  $2\sigma^2 p$ , which justifies Mallows'  $C_p$ .