

STAT 5630, Fall 2019

Neural Networks

Xiwei Tang, Ph.D. <xt4yj@virginia.edu>

University of Virginia

November 19, 2019

- Motivation
- Feedforward neural network
- Connections with other models
- Deep Neural Networks



Dogs vs. Cats classification problem at Kaggle: [Link](#)

- Neural Networks were first developed as models for the human brain, where each unit represents a neuron.
- The neurons fire when the total signal passed to that unit exceeds a certain threshold.
- The collective signal from all neurons tells you whether its a dog or a cat




FIGURE 11.2. *Schematic of a single hidden layer, feed-forward neural network.*

Formulate the problem

- Given a training set $\{x_i, y_i\}_{i=1}^n$,
 - For regression: $y_i \in \mathbb{R}^K$ is a K dimensional continuous outcome
 - For classification: $y_i \in \{1, 2, \dots, K\}$
- The goal is still to model the relationship

$$E(Y|X) = f(X)$$

- Instead of modeling the probabilities directly using X , we build M hidden neurons as a hidden layer between X and Y :

$$\begin{aligned}\mathbf{Z} &= (1, Z_1, Z_2, \dots, Z_M) \\ &= (1, \sigma(\mathbf{X}^\top \boldsymbol{\alpha}_1), \sigma(\mathbf{X}^\top \boldsymbol{\alpha}_2), \dots, \sigma(\mathbf{X}^\top \boldsymbol{\alpha}_M))\end{aligned}$$

- Where $\sigma(\cdot)$ is an activation function. Some examples?
- Then we model Y using the hidden layer variables \mathbf{Z} through some link function $g(\cdot)$

$$X \xrightarrow{\sigma(\cdot)} Z \xrightarrow{g(\cdot)} Y$$

Formulate the problem

- In **classification problems** (K class), we can use g_k to model the probability of $Y = k$, for $k = 1, \dots, K$:

$$g_k(\mathbf{Z}) = \frac{\exp(\mathbf{Z}^\top \beta_k)}{\sum_{l=1}^K \exp(\mathbf{Z}^\top \beta_l)}$$

- In **regression problems** (could be multidimensional), we can simply use a linear function to model the k th entry of Y :

$$g_k(\mathbf{Z}) = \mathbf{Z}^\top \beta_k$$

- The multidimensional function $\mathbf{f}(x)$ can be represented as a convoluted way of mapping $x \in \mathbb{R}^p$ to $y \in \mathbb{R}^K$

$$\mathbf{f}(x) = \mathbf{g} \circ \sigma(x)$$

- The notations \mathbf{g} and σ here are multidimensional.
- The parameters involved are: $\alpha_1, \dots, \alpha_M$, and β_1, \dots, β_K .

Examples of activation functions

- The activation function $\sigma(\cdot)$ takes a linear combination of the input variables, and output a scalar through nonlinear transformation. Examples:

- sigmoid:

$$\sigma(v) = \frac{1}{1 + e^{-v}} = \frac{e^v}{e^v + 1}$$

- hyperbolic tangent (tanh):

$$\sigma(v) = \frac{e^v - e^{-v}}{e^v + e^{-v}}$$

- rectified linear unit (ReLU):

$$\sigma(v) = \max(0, v), \quad \text{soft approx.} \quad \ln(1 + e^v)$$

- And many others: exponential linear unit, arctangent, etc.




FIGURE 11.3. Plot of the sigmoid function $\sigma(v) = 1/(1 + \exp(-v))$ (red curve), commonly used in the hidden layer of a neural network. Included are $\sigma(sv)$ for $s = \frac{1}{2}$ (blue curve) and $s = 10$ (purple curve). The scale parameter s controls the activation rate, and we can see that large s amounts to a hard activation at $v = 0$. Note that $\sigma(s(v - v_0))$ shifts the activation threshold from 0 to v_0 .

Formulate the problem

- Originally, a step function $\sigma(v)$ was considered as the activation function (to mimic the biological interpretation). Hence for each neuron, signal is triggered only when $x^T\alpha$ is above a certain threshold
- It was later recognized that the step function is not smooth enough for optimization, hence was replaced by a smoother threshold function, the sigmoid function
- “Feedforward” as signals can only pass to the next layer. There is no “cycle” in the model

- Try this a really cool website: <http://playground.tensorflow.org/>
- Implementation in R:
 - packages: `neuralnet`, `nnet`
 - `nnet` fits a single layer of hidden neurons; `neuralnet` can fit multiple layers
 - The initial parameters α 's and β 's are generated randomly and then optimized. The model fitting can be different depends on the initial value. To fix initial parameters: `nnet`: `Wts`; `neuralnet`: `startweights`
 - Number of neurons: `nnet`: `size`; `neuralnet`: `hidden` (if `hidden` is specified as a vector, then there will be multiple layers)

Universal Approximation Theorem (Cybenko, 1989; Hornik 1991)

Any continuous function $f(x)$ on the space $[0, 1]^p$ can be approximated (with any $\epsilon > 0$) by a finite set of neurons with a bounded monotone-increasing activation function $\varphi(\cdot)$:

$$\left| f(x) - \sum_{i=1}^n v_i \varphi(w_i^T x + b_i) \right| < \epsilon$$

for some v_i , w_i , and b_i . Hence, the functions defined by the neurons is dense.

- The parameters (weights) α 's and β 's need to be optimized.
- For a single hidden layer NN, we have

$$\{\alpha_1, \dots, \alpha_M\} : \quad M(p+1) \text{ weights}$$

$$\{\beta_1, \dots, \beta_K\} : \quad K(M+1) \text{ weights}$$

- where p is the number of non-intercept X features; M is the number of hidden neurons in a single layer; and K is the number of categories for classification.
- $K = 1$ if its a univariate regression problem.

- Neural Networks training is based on error minimization using a **Gradient Descent** algorithm, known as error **back-propagation**.
- For K classification, we minimize Deviance:

$$- \sum_{i=1}^n \sum_k^K \mathbf{1}\{y_i = k\} \log f_k(x_i)$$

- For univariate regression, we minimize RSS (since g is linear):

$$\sum_{i=1}^n (y_i - f(x_i))^2 = \sum_{i=1}^n (y_i - \beta_0 - \beta_1 \sigma(x^\top \alpha_1) - \cdots \sigma(x^\top \alpha_M))^2$$

Fitting Neural Networks

- The objective function can be written as

$$R(\boldsymbol{\theta}) = \sum_{i=1}^n R_i(\boldsymbol{\theta})$$



where R_i represents the deviance or residual sum of squares for the i th data point, and $\boldsymbol{\theta}$ represents an aggregated vector of all weights

- Initiate weights $\boldsymbol{\theta}^{(0)}$
- We then calculate the derivative wrt each of the weights evaluated at the current iteration value $\boldsymbol{\theta}^{(t)}$:

$$\sum_{i=1}^n \frac{\partial R_i}{\partial \beta_{km}} \bigg|_{\boldsymbol{\theta}=\boldsymbol{\theta}^{(t)}} \quad \sum_{i=1}^n \frac{\partial R_i}{\partial \alpha_{mj}} \bigg|_{\boldsymbol{\theta}=\boldsymbol{\theta}^{(t)}}$$

- Stochastic GD**: the summation can be taken over a **random subset** of the n samples

Gradient Descent vs. Stochastic Gradient Descent



Fitting Neural Networks

- The derivatives for $K = 1$ regression case is essentially

$$\frac{\partial R_i}{\beta_m} = -2(y_i - f(x_i))z_{mi}$$
$$\frac{\partial R_i}{\alpha_{ml}} = -2(y_i - f(x_i))\beta_m \sigma'(\alpha_m^T x_i) x_{il}$$

- Some redundant calculations can be saved in the above equations. The property is called **back-propagation**.
- We then do the update, at the t -th iteration

$$\beta_m^{(t+1)} = \beta_m^{(t)} - \gamma \sum_{i=1}^n \frac{\partial R_i}{\beta_m^{(t)}}$$
$$\alpha_{ml}^{(t+1)} = \alpha_{ml}^{(t)} - \gamma \sum_{i=1}^n \frac{\partial R_i}{\alpha_{ml}^{(t)}}$$

where γ is a step size for gradient descent.

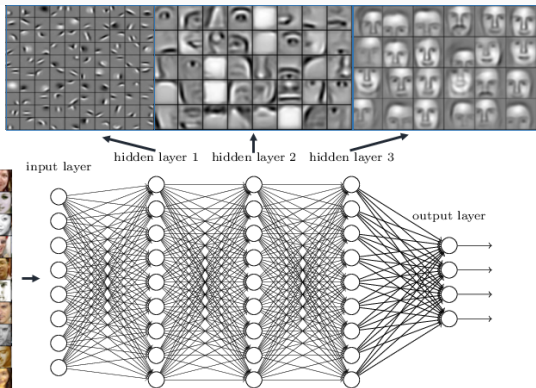
Fitting Neural Networks

- The derivatives can be calculated by Chain Rules
- The algorithm can be implemented by a forward-backward sweep over the network
- In the **forward** pass, compute the hidden variables and the output $\hat{f}(x_i)$ based on the current weights $\theta^{(t)}$
- In the **backward** pass, compute the derivatives, and update $\theta^{(t)} \rightarrow \theta^{(t+1)}$

- **Deep Neural Networks** are one type of deep learning models.
- Deep neural Networks are just ... Neural Networks with more than one hidden layer.
- But neural networks have been around for more than 70 years... why it gets popular just in recent years?
 - computational issues
 - a better way to generate/construct features
 - ...


Deep Neural Networks

Deep neural networks learn hierarchical feature representations



- One example is the Convolutional Neural Networks, which attempts to generate better features
- Instead of using all input features to create the linear combination, a “convolutional layer” builds neurons that each takes a subset (a local region) of the input features.
- This is motivated by the fact that biologically, the neurons only take signals from neighboring neurons.

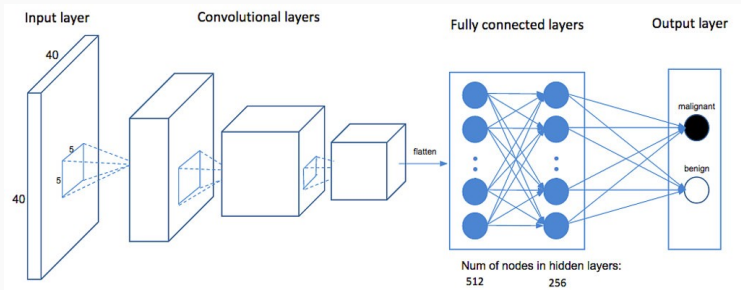
Deep Neural Networks: Feature Hierarchy



Hierarchical information processing in the brain


(Source: Simon Thorpe)

Convolutional Neural Networks



See this hand digit writing recognition [example](#), and this interesting application by [Tesla](#).

Convolutional Neural Networks



- The CNN consists of input an layer, convolution layers, pooling layers, a fully connected network and an output layer.
- Also known as shift invariant or space invariant artificial neural networks (SIANN).

Convolutional Layers

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

image 5*5

1	0	1
0	1	0
1	0	1

bias=0

filter 3*3

feature map 3*3

- Choose a filter which is assigned with certain weights $w_{m,n}$'s.

Convolutional Layers

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

image 5*5

1	0	1
0	1	0
1	0	1

bias=0

filter 3*3

4		

feature map 3*3

- Calculate the convolution (cross-correlation) by filter locally.

$$a_{i,j} = f\left(\sum_{m=0}^2 \sum_{n=0}^2 w_{m,n} x_{i+m,j+n} + w_b\right)$$

- $f(\cdot)$ is the activation function.

Convolutional Layers

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

image 5*5

1	0	1
0	1	0
1	0	1

bias=0


filter 3*3

4	3	

feature map 3*3

- Moving the filter across the image to obtain a feature map.
- Choose different stride size.
- Use multiple filters which results in multiple feature maps (depth)

Pooling Layers



- Dimension Reduction
- Choose pooling window size and stride size
- Max pooling, Mean Pooling, Median Pooling...
- Pooling on each layers individually

Some Points

- Vanishing gradient and exploding gradient
- Epoch: one epoch is when an ENTIRE dataset is passed forward and backward through the neural network only ONCE.
- Learning rate: gradient descent step size
- Batch size and number of batches (iterations)
- Momentum and normalization
- Weight Decay, regularization and dropout: overfitting

- Recurrent Neural Network (RNN)
- Long Short Term Memory Network, LSTM
- Recursive Neural Network, RNN

