

Reinforcement Learning

Lecture 14

So far... Supervised Learning

Data: (\mathbf{x}, \mathbf{y}) : \mathbf{x} is data, \mathbf{y} is label

Goal: Learn a function to map $\mathbf{x} \rightarrow \mathbf{y}$

Examples: Classification, regression, object detection, image captioning, etc.




→ A Cat

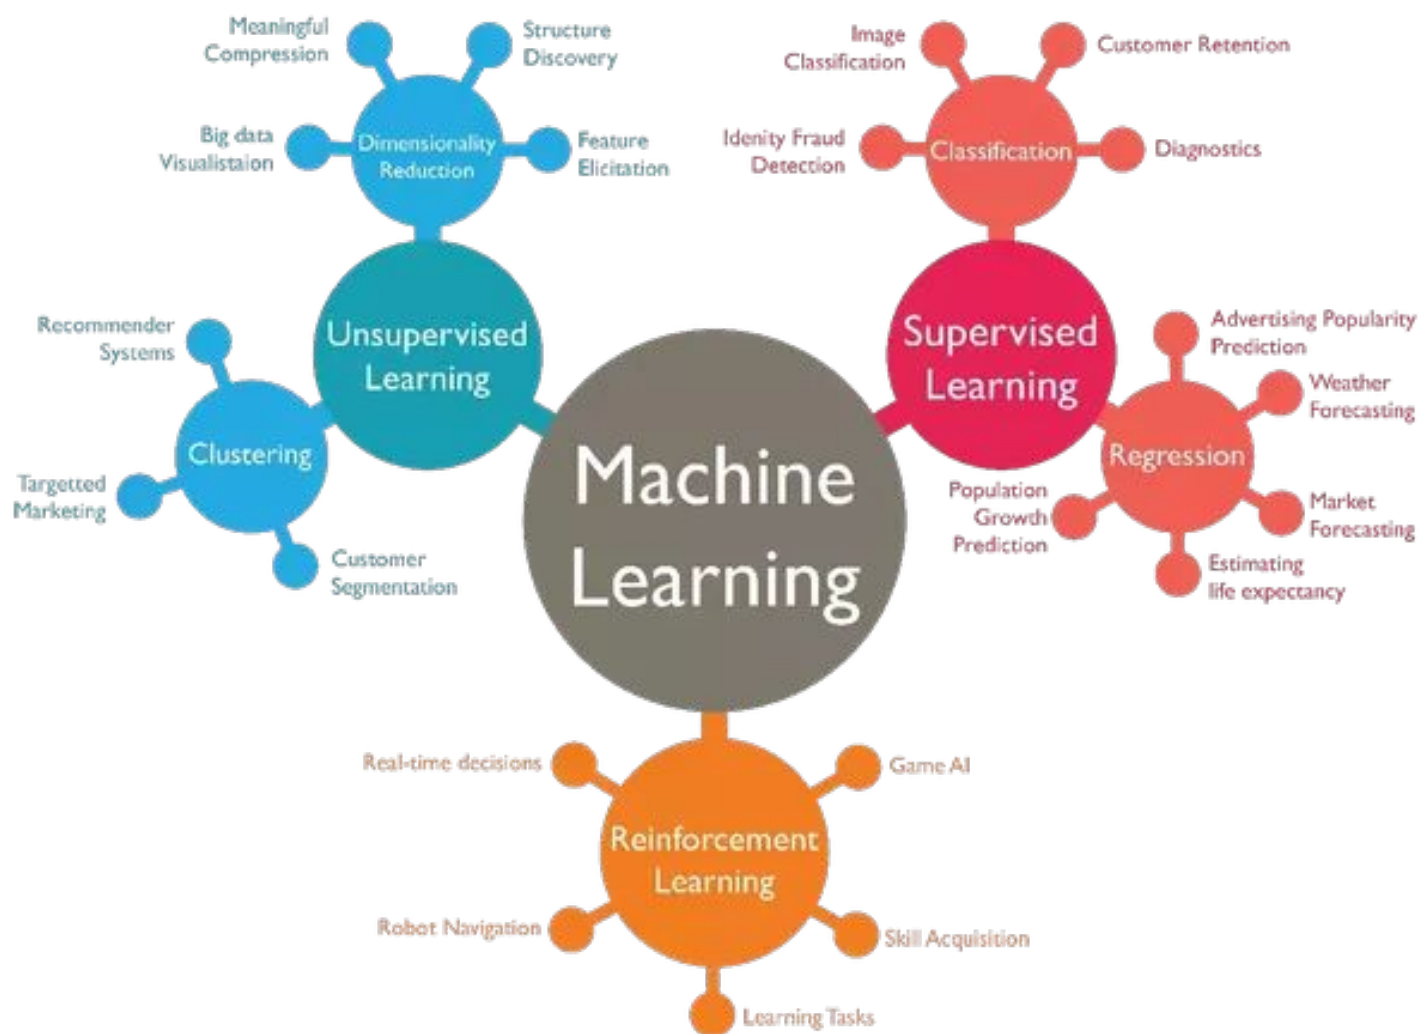
So far... Unsupervised Learning

Data: x Just data, no labels!

Goal: Learn some underlying hidden structure of the data

Examples: Clustering, dimensionality reduction, feature learning, etc.





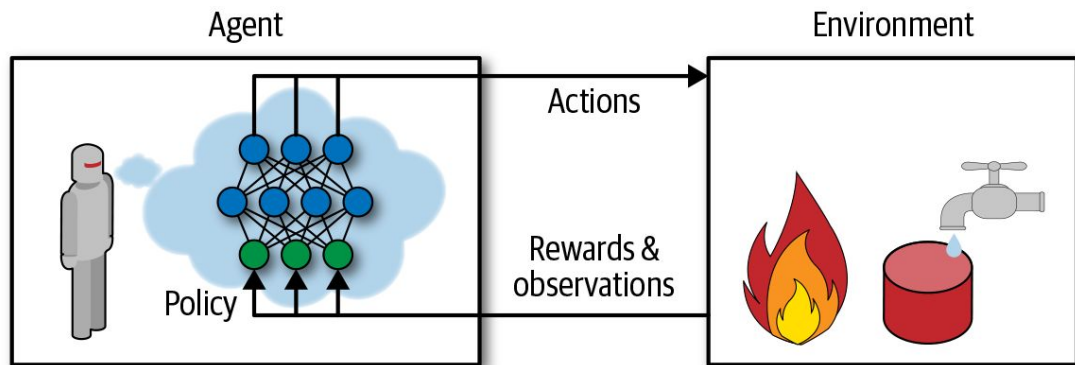
Today: Reinforcement Learning (RL)

Problems involving an **agent** interacting with an **environment**, which provides numeric reward signals.

At each step, the agent:

- Executes an **action**
- Observe a new **state**
- Receive some **reward**

Goal: Learn how to take actions from a policy in order to maximize reward




Example: Grasping Objects Problem

Goal: Pick an Object with different shape

State: Raw pixels from camera

Actions: Move arm, grasp

Reward: positive when pickup is successful




Example: Cart-Pole Balancing Problem

Goal: Balance the pole of top of a moving cart

State: Pole angle, angular speed, cart position, horizontal velocity

Actions: Horizontal force to the cart

Reward: 1 at each time step if the pole is upright



Example: DeepTraffic

Goal: train an RL agent that can successfully navigate through traffic

State: as a grid, where each cell will be the speed of the vehicle inside it

Actions: Accelerate, Break, Left, Right, No action

Reward: positive when high speed is maintained.



Example: College Life

Goal: Survival? Happiness?

State: Sight, hearing, taste, smell, touch, feel

Action: Think, move, speak

Reward: Grades? Money? Love?



PICK TWO, AND ONLY TWO.

Environment and Actions

Fully Observable (Chess) vs. **Partially Observable** (Poker)

Single Agent (Atari) vs. **Multi Agent** (DeepTraffic)

Deterministic (Cart Pole) vs. **Stochastic** (DeepTraffic)

Static (Chess) vs. **Dynamic** (DeepTraffic)

Discrete (Chess) vs. **Continuous** (Cart Pole)

RL in Humans

Humans appear to learn to walk through “very few examples” of trial and error.

“**How** we learn how to walk” is an open question...some possible answers:



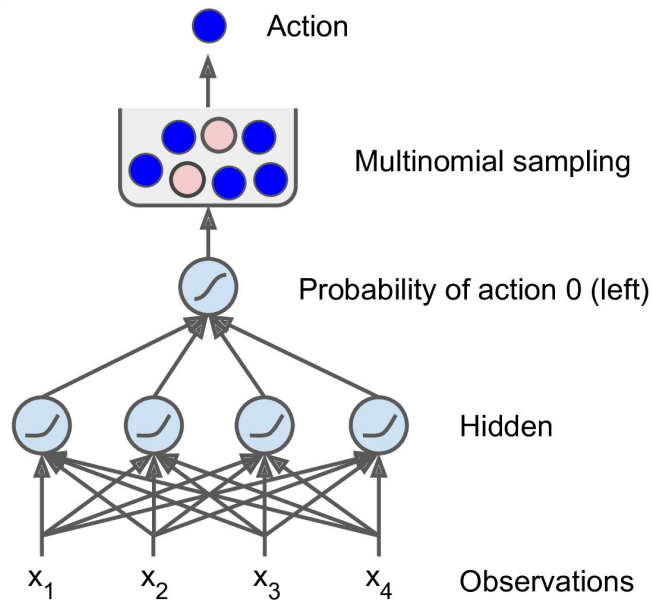
- **Hardware:** 230M years of bipedal movement data
- **Imitation Learning:** Observation of other human walking
- **Algorithms:** probably better than backprop and SGD

Deep RL


Deep? Deep RL = RL + Neural Networks

How? Trial and Error in a world that provides occasional rewards

⇒ a framework for learning to solve sequential decision-making problems.


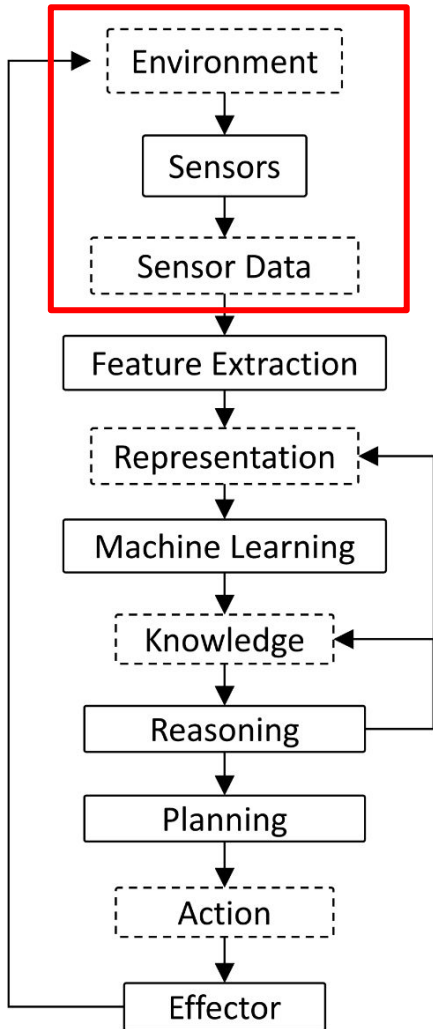


The Continuous Learning Cycle



What can be learned?

Sensors



Lidar



Camera
(Visible, Infrared)



Radar



GPS




Stereo Camera



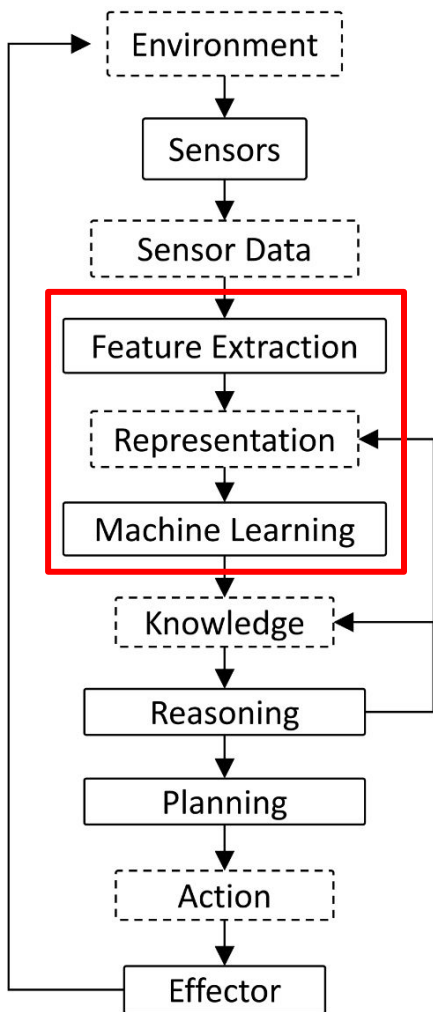
Microphone



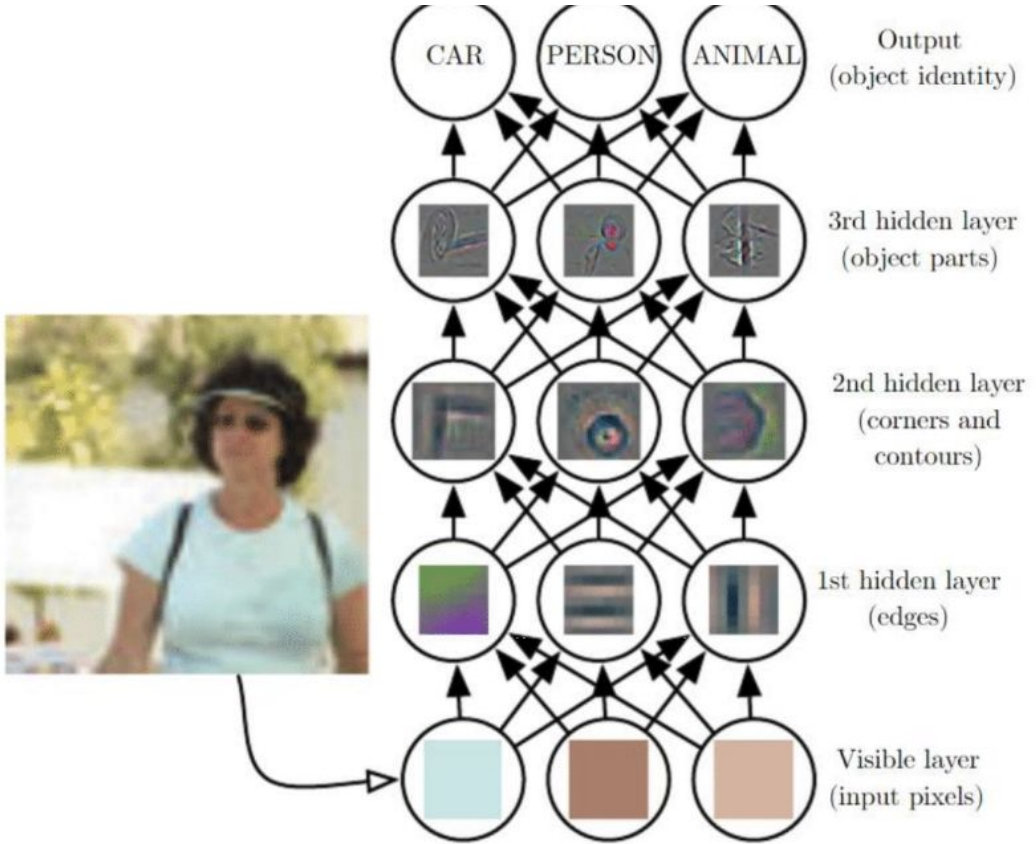
Networking
(Wired, Wireless)

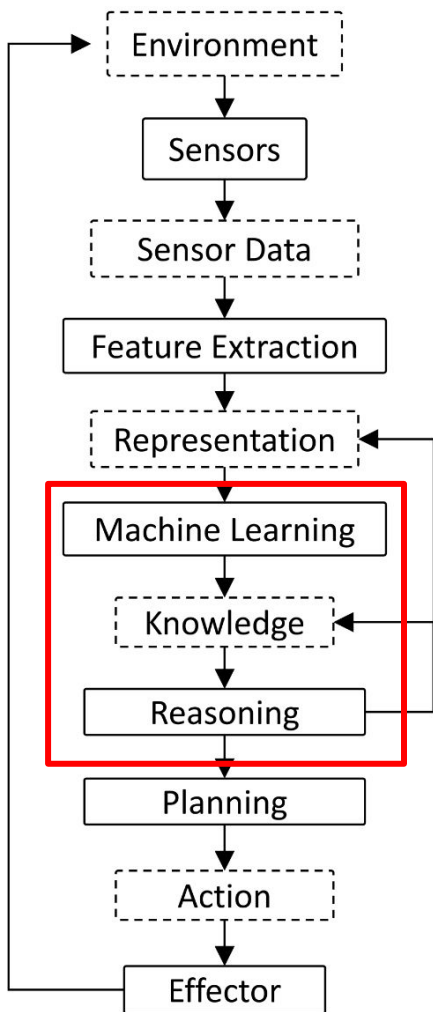


IMU



Representations






Knowledge / Reasoning

Image Recognition:
If it looks like a duck




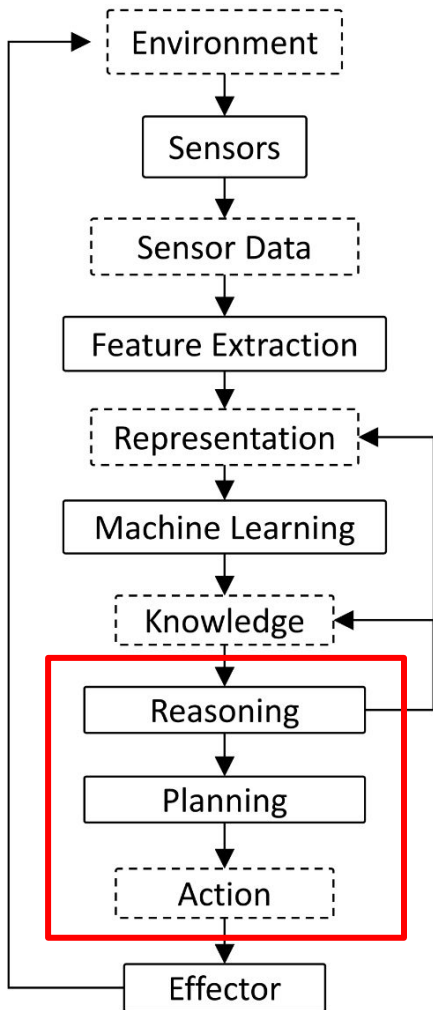
Audio Recognition:
Quacks like a duck



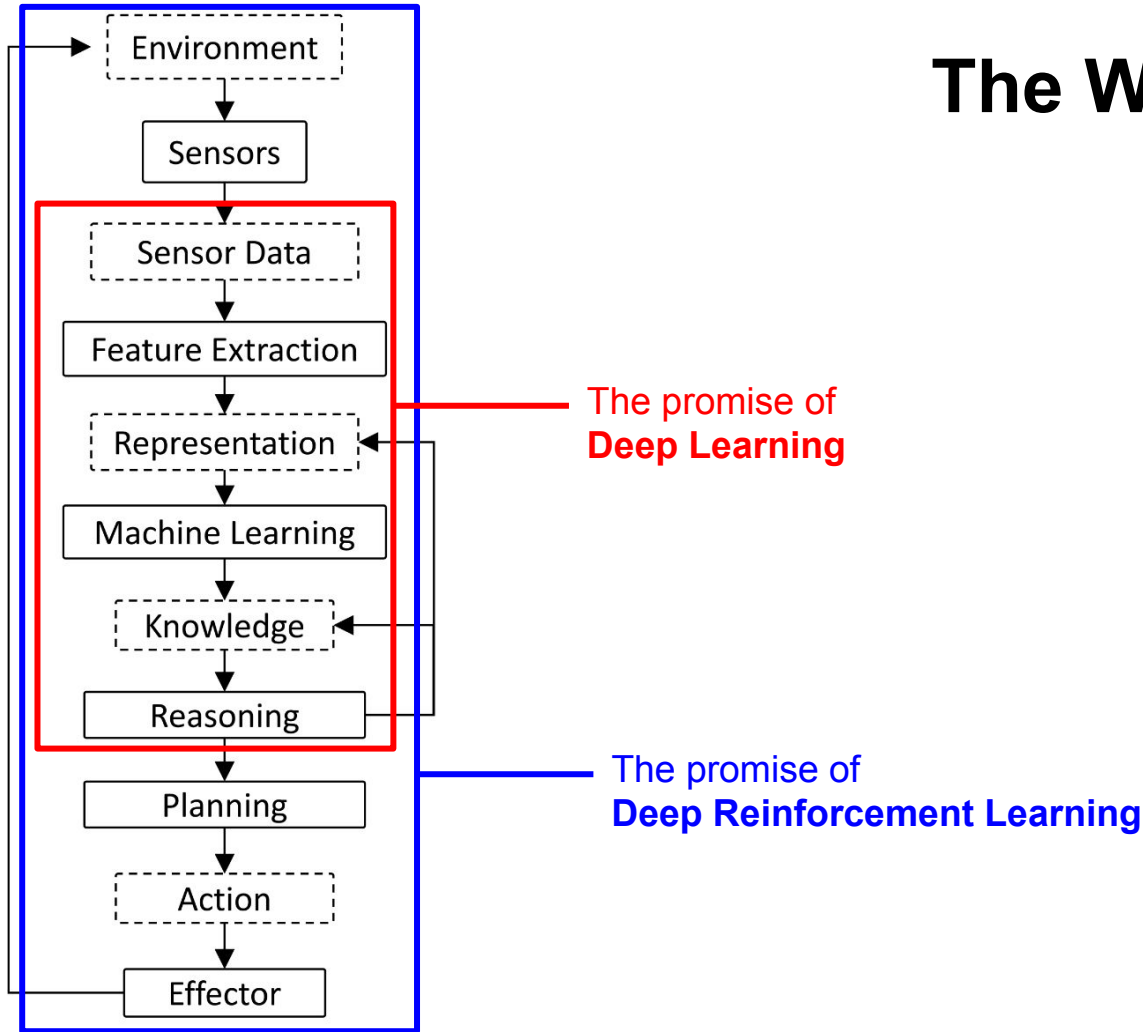
Activity Recognition:
Swims like a duck



Actions







The Whole Cycle



RL Applications

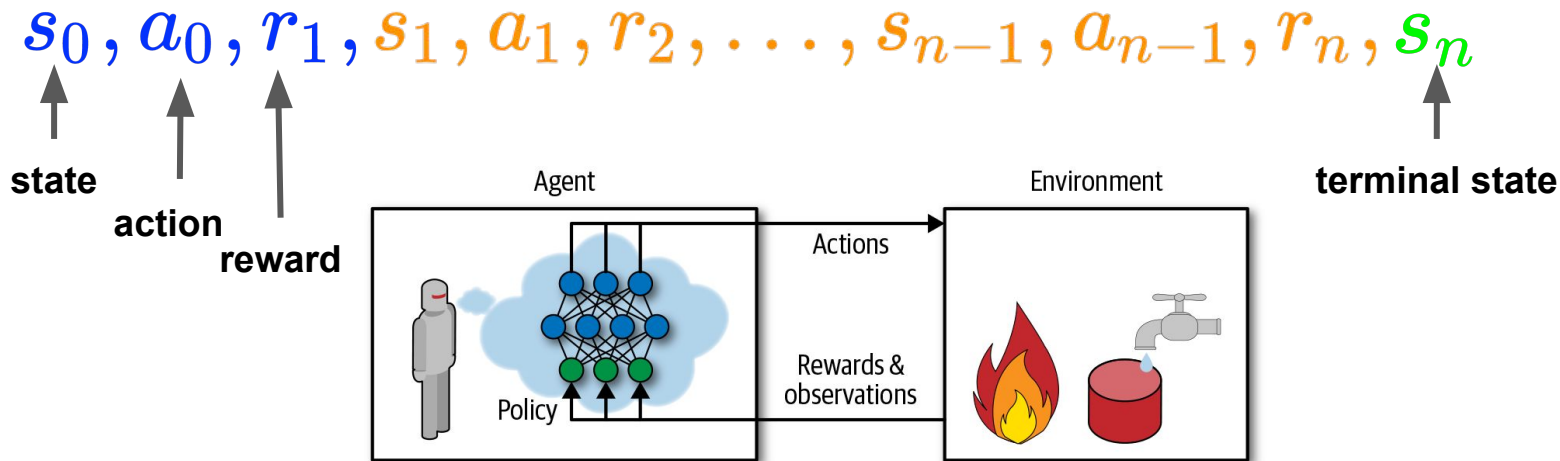
- a. Robotics
- b. Ms. Pac-man
- c. Go player
- d. Thermostat
- e. Automatic Trader



Major Components of an RL Agent

An RL agent may be directly or indirectly trying to learn an:


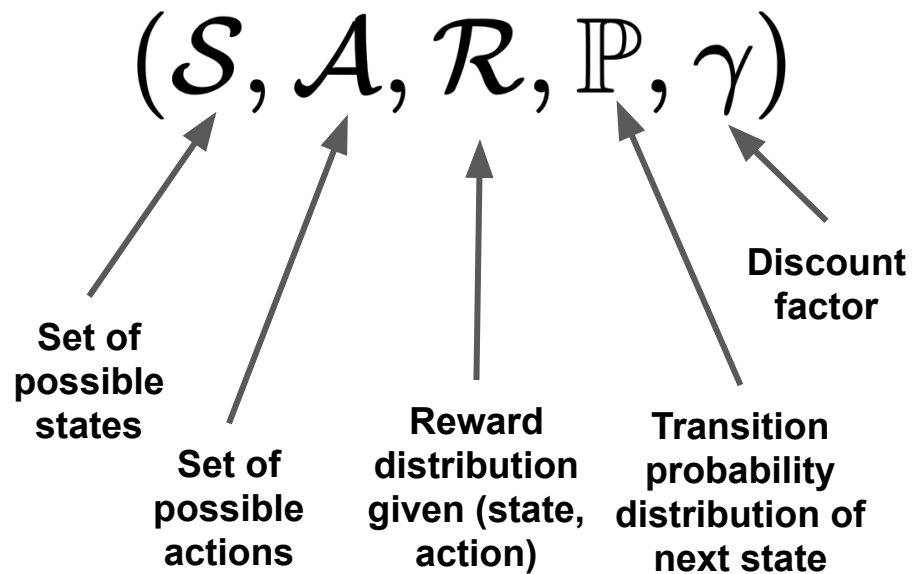
- **Policy:** agent's behavior function
- **Value Function:** how good is each state and/or action
- **Model:** agent's representation of the environment



Markov Decision Process

Markov Decision Process is the mathematical formulation for RL problem

Markov property: current state completely characterizes the state of the world.



Markov Decision Process

At time step $t = 0$, environment samples initial state $s_0 \sim \mathbb{P}(s_0)$

Then for $t = 0$ until done:

- Agent selects action a_t
- Environment samples reward $r_t \sim \mathcal{R}(\cdot | s_t, a_t)$
- Environment samples next state $s_{t+1} \sim \mathbb{P}(\cdot | s_t, a_t)$
- Agent receives the reward r_t and next state s_{t+1}

A policy π is a function from S to A that specifies what action to take in each state

Objective: find policy π^* that maximizes cumulative discounted reward

Maximize Reward

Future Reward: $R_t = r_t + r_{t+1} + r_{t+2} + \dots + r_n$

Discounted Future Reward: $R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^{n-t} r_n$

A good strategy for an agent would be to always choose an action that maximizes the discounted future reward

Why?

Uncertainty due the environment, partial observability

Real life example: Either Live it up today, or save \$ for tomorrow?

Moving in a grid world


Objective: reach one of the terminal states (stars) using the least number of actions.

actions = {

1. right 


2. left 

3. up 

4. down 

}


states




Set a negative “reward”
for each transition
(e.g. $r = -1$)

Policy to move in a grid world

Objective: reach one of the terminal states (stars) using the least number of actions.



Random Policy



Optimal Policy

3 Types of Reinforcement Learning

- **Model-based**: Learn the model then use and update it often.
- **Value-based**: Learn the state or state-action value, act by choosing best action, and explore if necessary
- **Policy-based**: Learn the stochastic policy function that maps state to action, act by sampling that policy.



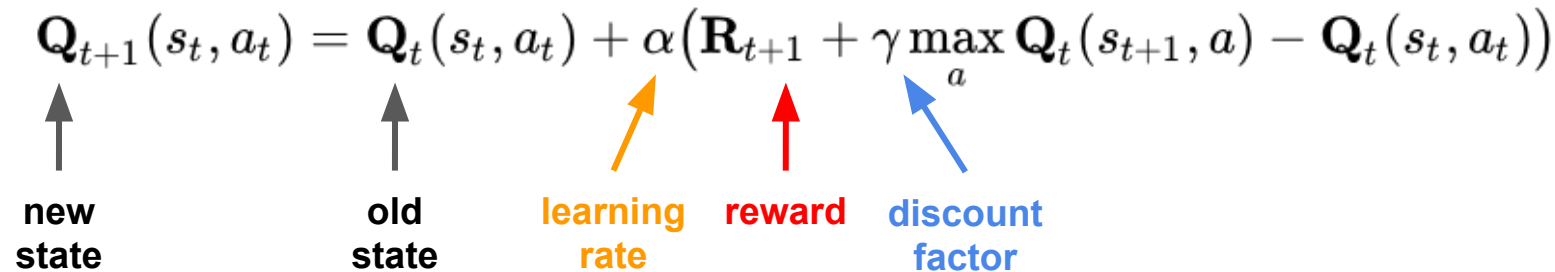
Value-based Method: Q-learning

Solving the Optimal Policy: Q-Learning

State-action value function: $Q_{\pi}(s,a)$ expected return when starting in s , performing a , and following π

Q-Learning: Use any policy to estimate Q that maximizes future reward:

- Q directly approximates Q^* (Bellman optimality equation)
- Independent of the policy being followed
- Only requirement: keep updating each (s,a) pair

$$Q_{t+1}(s_t, a_t) = Q_t(s_t, a_t) + \alpha (\mathbf{R}_{t+1} + \gamma \max_a Q_t(s_{t+1}, a) - Q_t(s_t, a_t))$$


The diagram shows the Q-learning update equation with arrows pointing to its components:

- new state**: points to $Q_{t+1}(s_t, a_t)$
- old state**: points to $Q_t(s_t, a_t)$
- learning rate**: points to α
- reward**: points to \mathbf{R}_{t+1}
- discount factor**: points to γ

Q-Learning: Value Iteration

$$Q_{t+1}(s_t, a_t) = Q_t(s_t, a_t) + \alpha(R_{t+1} + \gamma \max_a Q_t(s_{t+1}, a) - Q_t(s_t, a_t))$$


↑ ↑ ↗ ↑ ↖
new state old state learning rate reward discount factor

	A1	A2	A3	A4
S1	+1	+2	-1	0
S2	+2	0	+1	-2
S3	-1	+1	0	-2
S4	-2	0	+1	+1





```
initialize Q[num_states,num_actions] arbitrarily
observe initial state s
repeat
    select and carry out an action a
    observe reward r and new state s'
    Q[s,a] = Q[s,a] + α(r + γ maxa' Q[s',a'] - Q[s,a])
    s = s'
until terminated
```

Example: Moving in a grid world



Objective: reach one of the terminal states (stars) using the least number of actions.



actions = {

1. right 
2. left 
3. up 
4. down 

}



Optimal Policy


Exploration vs. Exploitation

Deterministic/greedy policy won't explore all actions

- Don't know anything about the environment at the beginning
- Need to try all actions to find the optimal one

ϵ -greedy policy

- With probability $1-\epsilon$ perform the greedy action, otherwise random action
- Slowly move toward greedy policy: $\epsilon \rightarrow 0$



Exploration vs. Exploitation Examples

- **Restaurant Selection**

- Exploitation: Go to your favourite restaurant
- Exploration: Try a new restaurant Online

- **Banner Ads**


- Exploitation: Show the most successful ads
- Exploration: Show a different ads

- **Oil Drilling**

- Exploitation: Drill at the best known location
- Exploration: Drill at a new location

- **Game Playing**

- Exploitation: Play the move you believe is best
- Exploration: Play an experimental move



Q-Learning: Representation Matters

Unfortunately, value iteration is **impractical**

- Limited states/actions
- Cannot generalize to unobserved states

Think about the **Breakout** Arcade game

State: screen pixels

- Image size: 84 x 84 (resized)
- Consecutive 4 images
- Grayscale with 256 gray levels
- → **256^{84x84x4}** rows in the Q-table! ($256^{28,224} = 10^{69,970} \gg 10^{82}$ atoms in the universe)




Deep RL = RL + Neural Networks


Use a deep neural network to approximate
Q-function → Deep Q-Network (DQN):

$$Q(s, a; \theta) \approx Q^*(s, a)$$

↑
network
parameters



Deep Q-Network (DQN) Architecture



Layer	Input	Filter size	Stride	Num filters	Activation	Output
conv1	84x84x4	8x8	4	32	ReLU	20x20x32
conv2	20x20x32	4x4	2	64	ReLU	9x9x64
conv3	9x9x64	3x3	1	64	ReLU	7x7x64
fc4	7x7x64			512	ReLU	512
fc5	512			18	Linear	18


Demo video



Experience Replay

- Current Q-network parameters determines next training examples → can lead to **bad feedback loop**
- Stores experience (actions, state transitions, and rewards) and create **mini-batches** from them for the training process
- Continually update a **replay memory table** of transitions as game experience are played
- Update Q-network on random mini-batch of transitions from the replay memory, instead of consecutive samples

DQN on Atari



Policy-based Method: Policy Gradient

Policy Gradients

A problem with Q-learning is that the Q-function can be very complicated!


Example: a robot grasping an object has a **very high dimensional state** \Rightarrow hard to learn exact value of every (state, action) pair

The policy could be much simpler: just close your hand/claw


Can we optimize a policy **directly** by finding the best one from a collection of policies?



Policy Gradients -- on Pong




- 80 x 80 image (difference image)
- 2 actions: up or down
- 200,000 Pong games
- This is a step towards **general purpose AI** !



Policy Gradients -- on Pong



Action trajectories/**series**:



Policy Gradients (PG)

Formally, let's define a class of policies: $\Pi = \{\pi_{\theta}, \theta \in \mathbb{R}^m\}$

For each policy, define its value:

$$J(\theta) = \mathbb{E}\left[\sum_{t \geq 0} \gamma^t r_t \mid \pi_{\theta}\right]$$

We want to find the optimal policy: $\theta^* = \arg \max_{\theta} J(\theta)$

How to do this? \rightarrow **Gradient Descent** (Ascent) $\nabla_{\theta} J(\theta)$ on policy parameter θ !

Policy Gradients (PG)

Mathematically, we can rewrite $J()$ in terms of action trajectory:

$$J(\theta) = \mathbb{E}_{\tau \sim p(\tau; \theta)} [r(\tau)]$$

Where $\mathbf{r}(\tau)$ is the reward of a trajectory: $\tau = (s_0, a_0, r_0, s_1, a_1, r_2, \dots)$

Gradient Estimator (skipping the derivation...):

$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

Interpretation:

- If $\mathbf{r}(\tau)$ is high, push up the probabilities of the seen actions
- If $\mathbf{r}(\tau)$ is low, push down the probabilities of the seen actions

PG comparing to DQN


Pros:

- + **Messy World:** If Q function is too complex to learn, DQN may fail while PG will still learn a good policy
- + **Speed:** Faster convergence
- + **Stochastic:** PG is capable of learning stochastic policies while DQN cannot
- + **Continuous actions:** It's easier to model PG on continuous space

Cons:

- **Data:** Sample Inefficient (need more data)
- **Stability:** Less stable during training process
- **Credit assignment:** Poor assignment to (state, action) pairs for delayed rewards

The problem with Policy Gradients



We have to wait until the end of a trajectory to calculate the reward. If the reward were high, all actions that we took were good, even if some were **really bad**

As a consequence, we need to have a **LOT** of samples to have an optimal policy. This means slow learning and long time to converge

What if we can do update at each step?!

Hybrid Method: Actor-Critic

Introducing Actor-Critic Algorithm

Using **two** neural networks:

1. **An Actor** that measures how good the action taken (**value-based**)
2. **A Critic** that controls how our actor behaves (**policy-based**)





Actor Critic Algorithm

- The actor decides which action to take, and the critic tells the actor how good its action was and how it should adjust
- Alleviates the task of the critic as it only has to learn the values of (state, action) pairs generated by the policy
- Can also incorporate Q-learning tricks (e.g. experience replay)
- **Remark:** we can define by the **advantage function** how much an action was better than expected $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$
- Using this, our gradient estimator becomes:

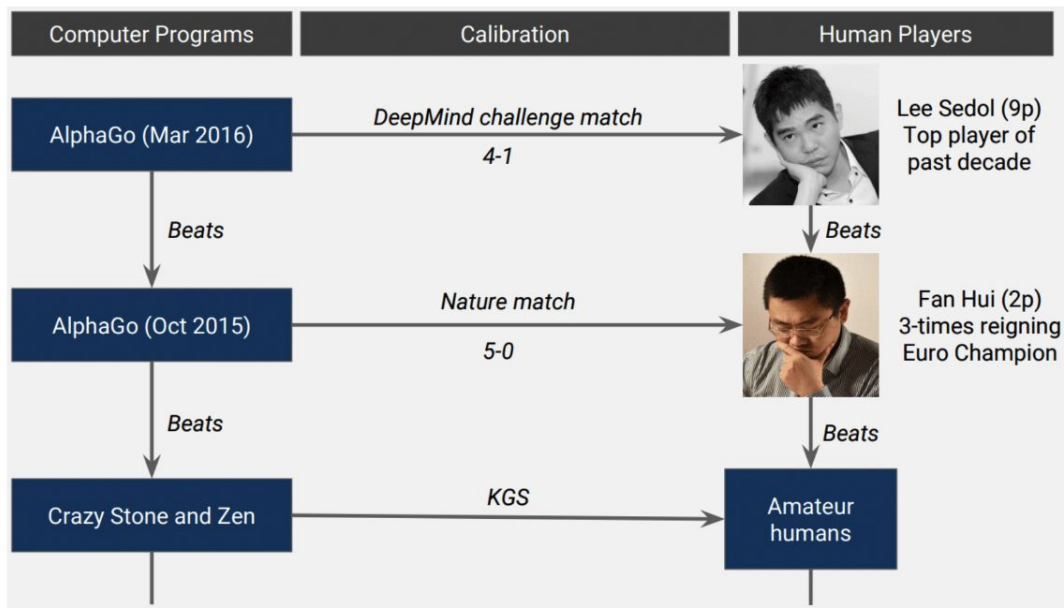

$$\nabla_\theta J(\theta) \approx \sum_{t \geq 0} (Q^{\pi_\theta}(s_t, a_t) - V^{\pi_\theta}(s_t)) \nabla_\theta \log \pi_\theta(a_t | s_t)$$

Application: Game of Go




Game size	Board size N	3^N	Percent legal	legal game positions (A094777) ^[11]
1×1	1	3	33%	1
2×2	4	81	70%	57
3×3	9	19,683	64%	12,675
4×4	16	43,046,721	56%	24,318,165
5×5	25	8.47×10^{11}	49%	4.1×10^{11}
9×9	81	4.4×10^{38}	23.4%	1.039×10^{38}
13×13	169	4.3×10^{80}	8.66%	$3.72497923 \times 10^{79}$
19×19	361	1.74×10^{172}	1.196%	$2.08168199382 \times 10^{170}$

Alpha Go



AlphaGo and variants



AlphaGo [Nature 2016]:

- Required many engineering tricks
- Bootstrapped from human play
- Beat 18-time world champion Lee Sedol

AlphaGo Zero [Nature 2017]:

- Simplified and elegant version of AlphaGo
- No longer bootstrapped from human play
- Beat (at the time) #1 world ranked Ke Jie

Alpha Zero: [Science 2018]


- Generalized to beat world champion programs on chess and shogi as well

Alpha Zero in action

Chess




AlphaZero vs. Stockfish



Shogi




AlphaZero vs. Elmo



Go



AlphaZero vs. AGO



AZ wins ■ AZ draws ■ AZ loses ■ AZ white ○ AZ black ●

Summary: Learning Objectives

- ✓ Overview of Reinforcement Learning (RL)
- ✓ Value-based Q-Learning method and Deep Q-Network
- ✓ A Policy-based method called Policy Gradients
- ✓ The Actor-critic algorithm and an application in AlphaGo

Next Steps in RL

Background

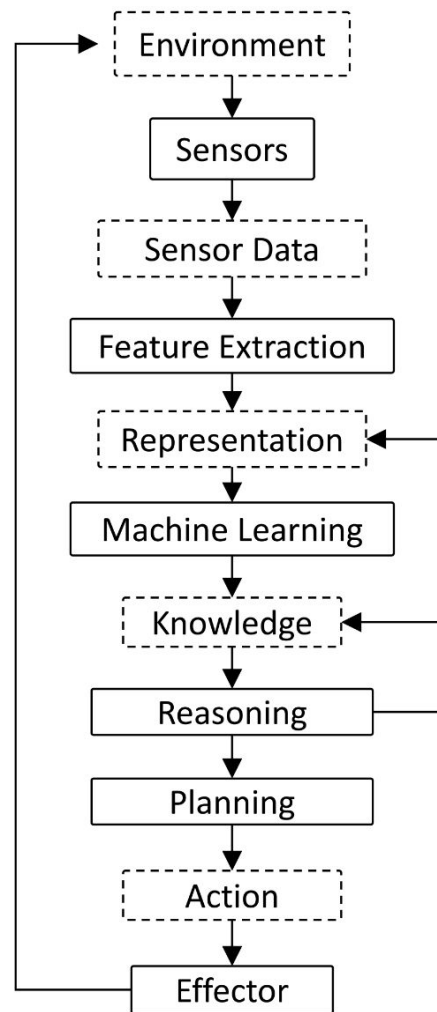
- Fundamentals in probability, statistics, multivariate calculus.
- Deep learning basics
- Deep RL basics
- TensorFlow (or PyTorch)

Learn by doing

- Implement core deep RL algorithms (discussed today)
- Look for tricks in papers that were key to get it to work
- Iterate fast in simple environments

Research

- Improve on an existing approach
- Focus on an unsolved task / benchmark
- Create a new task / problem that hasn't been addressed



Acknowledgement

Slides contain materials and figures reproduced from Lex Fridman at MIT and Serena Yeung at Stanford for educational purposes only.




Bonus Slides

PG Application: AlphaGo

Mix of supervised learning and reinforcement learning


How to beat the Go world champion:

- Featurize the board (stone color, move legality, bias,...)
- Initialize policy network with supervised training from professional go games, then continue training using policy gradient (**play against itself** from random previous iterations, +1 / -1 reward for winning / losing)
- Also learn value network (the critic)
- Finally, combine policy and value networks in a Monte Carlo Tree Search algorithm to select actions by look-ahead search




Alpha Zero in action

Elo rating is a method to calculate relative skill level of player in a zero-sum based game such as chess. Each training step represents 4,096 board positions



AlphaZero



DQN and double DQN

Loss function:

$$L = \mathbb{E}[(\underbrace{r + \gamma \max_{a'} \mathbf{Q}(s', a')}_{\text{Target (y)}} - \underbrace{\mathbf{Q}(s, a; \theta)}_{\text{prediction}})^2]$$

DQN: same network for both Q

Double DQN: separate network for each Q to help reduce bias introduced by the inaccuracies of Q network at the beginning of training