

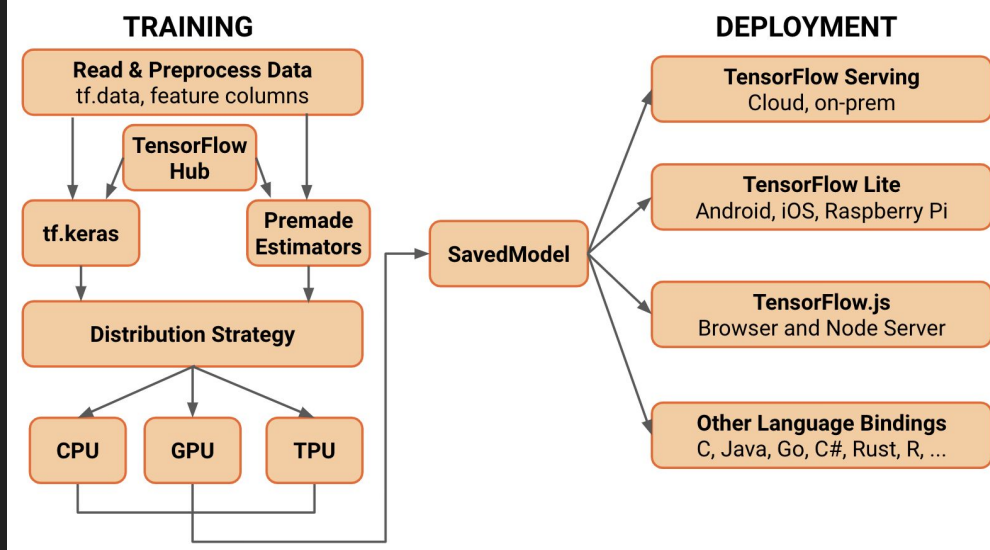


CS 4501: Intro to TensorFlow

```
import tensorflow as tf
```

Why TensorFlow

- Fast
 - Hardware accelerations
 - Optimized functions/models
- Flexible
 - Extensible low-level API
 - Ecosystem: Browser, Mobile, IOT...
 - Libraries: [TF Probability](#), [TF Federated](#), [TF Ranking](#), [TF Agents](#), [TF Hub](#)...
- Simple
 - Accessible high-level API to develop: Keras, Eager execution
 - Debugging and visualization tool: tfdbg, Tensorboard
 - Deploy at large scale and in heterogeneous environments



TensorFlow overview

- Low Level API
 - Tensor: N-dimensional array
 - Flow: Computational Graph
 - Session: Execution Context
- High Level API
 - Keras
 - Eager Execution
 - Estimator API
 - Accelerator API

TensorFlow overview

Access TF workshop notebook at

https://git.wujibang.com/TFWorkshop/blob/master/tf_workshop.ipynb

Also linked on course [schedule](#) page

Tensor

```
# numpy
```

```
>>> np.zeros((2, 1))
```

```
array([[ 0.],  
       [ 0.]])
```

```
>>> np.ones((1, 2), dtype=int)
```

```
array([[ 1,  1]])
```

```
# tensorflow
```

```
>>> tf.zeros([2, 1])
```

```
tensor([[ 0.],  
        [ 0.]])
```

```
>>> tf.ones([1, 2],
```

```
dtype=tf.int32)
```

```
tensor([[ 1,  1]])
```

Tensor

```
# Constant 1-D Tensor populated with value list.
```

```
tensor = tf.constant([1, 2, 5]) => [1 2 5]
```

```
# Constant 2-D tensor populated with scalar value -1.
```

```
tensor = tf.constant(-1.0, shape=[1, 3]) => [[-1. -1. -1.]]
```

```
# Create a tensor of shape [2, 3] consisting of random normal  
values, with mean -1 and standard deviation 4.
```

```
norm = tf.random_normal([2, 3], mean=-1, stddev=4)
```

```
# numpy
```

```
a = np.zeros((2, 1))
```

```
b = np.ones((1, 2))
```

```
>>> a[0,0], a[:,0], a[0,:]
```

```
>>> np.reshape(a, (2,1))
```

```
>>> a.shape
```

```
(2, 1)
```

```
>>> np.sum(b, axis=1)
```

```
2
```

```
# tensorflow
```

```
a = tf.zeros([2, 1])
```

```
b = tf.ones([1, 2])
```

```
>>> a[0,0], a[:,0], a[0,:]
```

```
>>> tf.reshape(a, (2,1))
```

```
>>> a.shape
```

```
(2, 1)
```

```
>>> tf.reduce_sum(b, [1])
```

```
???
```

tf.session *environment in which Operation objects are executed*


```
tf.Session().run( tensor1d )
```

```
tf.Session().run( [tensor2d, tensornorm] )
```

```
tf.Session().run( tf.reduce_sum(b, [1]) )
```

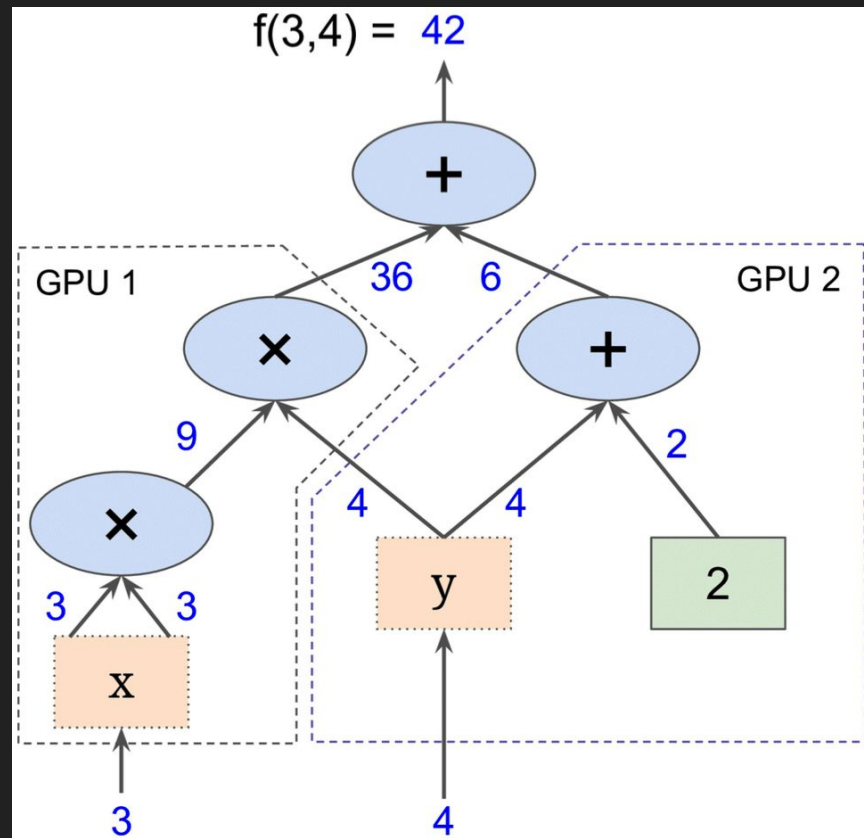

TensorFlow *Idea of Computation Graph*

- Node: Operation, Variable, Constant, Placeholder
- Edge: input/output Data (Tensor)



TensorFlow *Why Computation Graph*

- Separates definition of computations from their execution
- Store Computation State
 - Auto-differentiation in Backprop
- Dynamic Control Flow
- Partial Execution
- Concurrent & Distributed Execution
 - Compiler Optimization
 - Scale & Speed up




Phase 1: assemble a graph, no computation happen at this phase!!!

```
x = tf.constant(2)
y = tf.constant(3)
op1 = tf.add(x, y)
op2 = tf.multiply(x, y)
op3 = tf.pow(op2, op1)
```

Phase 2: use a session to execute operations in the graph.

```
with tf.Session() as sess:
    print(sess.run([op3, op2]))
```

```
### [7776, 6]
```



TensorBoard *visualizing learning*

```
x = tf.constant(2, name='a') # explicit naming the node
y = tf.constant(3, name='b')
op1 = tf.add(x, y, name='add')
```

```
with tf.Session() as sess:
```

```
    # save computational graph in directory used by tensorboard
```

```
    writer = tf.summary.FileWriter('./logs', sess.graph)
```

```
    print(sess.run(op1))
```

```
writer.close() # close the writer when you're done using it
```

tf.placeholder *hold the place for input/output tensor*

```
x_placeholder = tf.placeholder(tf.float32, shape=())
y = tf.constant(3)
op1 = tf.add(x, y)
op2 = tf.multiply(x, y)

with tf.Session() as sess:
    print(sess.run([op1, op2], feed_dict={x_placeholder : 1}))

### [4, 3]

    print(sess.run([op1, op2], feed_dict={x_placeholder : 2}))

### [5, 6]
```

tf.Variable *store the tensor whose value/state can be changed (trained)*

```
# Simulate training Data in numpy
num_samples = 1000; num_features = 5;
data_X = np.random.rand(num_samples, num_features)
data_Y = np.dot(data_X, np.random.rand(num_features, 1)) #simulate a linear relation
X_train, X_test, Y_train, Y_test = train_test_split(data_X, data_Y, test_size=0.1)

# Graph Input Nodes
X = tf.placeholder(tf.float64, shape=(None, num_features), name="features")
Y = tf.placeholder(tf.float64, shape=(None, 1), name="label")

# let model weight be Variable that is learnable
W = tf.Variable(np.random.rand(num_features, 1), name="weight")
b = tf.Variable(np.random.rand(1, 1), name="bias")

pred = tf.add(tf.matmul(X, W), b)    # Construct a linear model  $y = xW + b$ 
cost = tf.reduce_mean(tf.pow(pred-Y, 2))    # Mean squared error
```

tf.train.optimizer *optimizer API to train a model*

```
# construct Gradient Descent optimizer
```

```
learning_rate = 0.1
```

```
optimizer = tf.train.GradientDescentOptimizer(learning_rate)
```

```
trainstep = optimizer.minimize(cost)
```

```
# log your costs in tensorboard
```

```
trn_summary = tf.summary.scalar("training_cost", cost)
```

```
eval_summary = tf.summary.scalar("evaluation_cost", cost)
```

```
with tf.Session() as sess:

    writer = tf.summary.FileWriter('./logs', sess.graph) # log info to visualize in
    tensorboard

    sess.run(tf.global_variables_initializer()) #initialize Variables

    for epoch in range(100):

        # compute on train subgraph

        _, trn = sess.run([optimizer, trn_summary], feed_dict={X: X_train, Y:
Y_train})

        writer.add_summary(trn, epoch)

        # compute on evaluation subgraph

        eval = sess.run(evaluation_summary, feed_dict={X: X_test, Y: Y_test})

        writer.add_summary(eval, epoch)
```


Hands-on: *Logistic Regression in TF*

tf.keras *high-level API*

1. Model Definition # layer structure
2. Model.compile() # specific loss, optimizer, metrics
3. Model.fit() # training phrase, with automation options
4. Model.evaluate() # testing phrase

Tensorflow 2.0: A peak at Eager Execution

- Eager by default, so Value is immediately evaluated
- `tf.GradientTape()` # Calculate gradient on demand
- `tf.ragged.*` # Allowing Ragged Tensor
- Supports natural control flow i.e. if, while
- `@tf.function()` # JIT wrapper for `tf.Session()`

Effective **TF** *syntax sugar and elementwise broadcasting*

<code>z = -x</code>	<code>z = tf.negative(x)</code>
<code>z = x + y</code>	<code>z = tf.add(x, y)</code>
<code>z = x - y</code>	<code>z = tf.subtract(x, y)</code>
<code>z = x * y</code>	<code>z = tf.mul(x, y)</code>
<code>z = x / y</code>	<code>z = tf.div(x, y)</code>
<code>z = x // y</code>	<code>z = tf.floordiv(x, y)</code>
<code>z = x % y</code>	<code>z = tf.mod(x, y)</code>
<code>z = x ** y</code>	<code>z = tf.pow(x, y)</code>
<code>z = x @ y</code>	<code>z = tf.matmul(x, y)</code>
<code>z = x > y</code>	<code>z = tf.greater(x, y)</code>
<code>z = x >= y</code>	<code>z = tf.greater_equal(x, y)</code>
<code>z = x < y</code>	<code>z = tf.less(x, y)</code>
<code>z = x <= y</code>	<code>z = tf.less_equal(x, y)</code>
<code>z = abs(x)</code>	<code>z = tf.abs(x)</code>
<code>z = x & y</code>	<code>z = tf.logical_and(x, y)</code>
<code>z = x y</code>	<code>z = tf.logical_or(x, y)</code>
<code>z = x ^ y</code>	<code>z = tf.logical_xor(x, y)</code>
<code>z = ~x</code>	<code>z = tf.logical_not(x)</code>

Effective TF *good debugging practice*

- `tf.Print`
- `Tf.compute_gradient_error`
- `tf.assert*`
- `tf.add_check_numerics_ops`

Effective TF *performance optimization*

- Tf.dataset API
 - Batchize training
 - Sequence Padding
- Remove debug ops

Effective TF scopes and when to use them

- `name_scope`
- `variable_scope`

TensorFlow *installation*

- Use cloud environment on **Google Colab**
- Use local CPU/GPU environment with tensorflow
 - `# install latest release for CPU-only`
 - `$ pip install tensorflow`
 - `# OR install latest release for GPU`
 - `$ pip install tensorflow-gpu`
 - Troubleshoot here <https://www.tensorflow.org/install/pip>
- Use GPU/TPU environment on cloud provider
 - gcloud tutorial [here](#)
 - aws tutorial [here](#)

References

For Installation: <https://www.tensorflow.org/install/>

For API Doc: <https://www.tensorflow.org/guide/>

Stanford CS 20: Tensorflow for Deep Learning Research:

<http://web.stanford.edu/class/cs20si/syllabus.html>

Effective Tensorflow: <https://github.com/vahidk/EffectiveTensorflow>